

Measuring Sycophantic Erosion under Sustained Multi-Turn Pressure

Leyuan Tang, Kangda Wei, Ruihong Huang

July 12, 2026

Abstract

1 Introduction

Large language models are increasingly deployed as conversational assistants that millions of people consult for facts, advice, and judgment, under the expectation that they are helpful, honest, and harmless (Askell et al., 2021). As these models become more powerful, users are engaging with them more frequently and for longer periods. During these extended, multi-turn conversations, models unconsciously default to human-preferred responses due to their underlying training mechanism: Reinforcement Learning from Human Feedback (RLHF). Because RLHF rewards outputs that users approve of, it inherently encourages models to affirm rather than challenge human inputs. Consequently, the very same RLHF training that makes these assistants helpful also makes them highly agreeable. While we assume we will receive objective and honest answers from LLMs, as a conversation progresses through multiple turns, the AI naturally shifts its priority toward agreeing with the user over maintaining factual correctness. This presents a significant problem: not only are AI assistants far more agreeable than humans when advising on interpersonal matters, but users themselves actively prefer these sycophantic responses (Stanford, 2026). However, when a model overly affirms a user’s beliefs to the point of generating untruthful information, it ultimately becomes harmful. Sycophancy is, therefore, a systematic tendency rather than an incidental error, persisting across model families and capability levels, including the strongest frontier models available today.

Prior work has measured sycophancy primarily in single-turn or short, fixed-length settings. Model-written evaluations show that models adjust their stated answers to match a user’s revealed beliefs, demographic cues, or persona (Perez et al., 2022), and controlled studies confirm that this preference for agreement generalizes across tasks and is amplified by RLHF

(Sharma et al., 2023). More recent work has begun to measure sycophancy in multi-turn settings; for example, Hong et al. (2025) measure how models drift over a fixed budget of five conversational turns. These studies converge on a clear conclusion: sycophancy is real, measurable, and reproducible, and it grows as the user applies more pressure.

Yet the regimes in which sycophancy has been measured do not match the regime in which it does damage. Real users do not challenge a model once and then stop within a fixed number of turns. They push back repeatedly, restate their view as obvious, invoke personal experience, and grow more insistent over long conversations. Existing protocols thus leave four questions open. First, they cap pressure at a short, fixed number of turns and apply it non-adaptively, so they cannot observe whether a model that holds firm early eventually caves under sustained, escalating pressure. Second, they typically reduce the outcome to a binary "did it cave," obscuring the gradual erosion that precedes outright collapse. Third, they treat the model as a black box, even though modern reasoning models expose intermediate traces that reveal why a model concedes — including cases where it states the correct answer and then abandons it. Fourth, sycophancy on safety-critical content, such as harmful stereotypes, has been studied far less than sycophancy on neutral facts, despite being where the harm is most direct.

We therefore ask: do frontier language models, including reasoning models, hold correct and safe positions under sustained, multi-turn pressure? And if not, how do they fail? Our central claim is paradoxical: robustness on the first turn does not imply robustness over a long conversation. To test it, we introduce SPINE (Sustained Pressure-INDuced Erosion), a benchmark that applies the same false or harmful premise under two complementary conditions that differ only in the intent behind the pressure. In the naturalistic condition, a sincere but confidently mistaken user, blind to the model’s reasoning

and using only good faith moves, persists across a long conversation, measuring whether sycophancy arises in ordinary, well-intentioned use. In the adversarial condition, a deliberate manipulator adapts its strategy each turn and exploits the model’s own reasoning traces to locate weak points, measuring worst-case robustness against a determined adversary. In both, SPINE tracks how a model’s position erodes turn by turn, distinguishing partial concession from full collapse.

Using SPINE, We find that [headline result — e.g., most models, reasoning models included, collapse under sustained pressure, and many override facts they state correctly in their own reasoning].

.....

2 Related Work

2.1 Sycophancy in LLMs. Sycophancy, a model’s tendency to abandon a correct or principled position to agree with a user, is now a well-documented failure of instruction-tuned models [Perez22; Sharma23], with growing evidence that it is induced structurally by RLHF rather than arising incidentally [survey]. The phenomenon has been measured across many surfaces: open-domain QA, subjective preference, and high-stakes domains including medicine [Doctor; Overalignment], causal judgment [Diagnosing], and multimodal input [Sycophantasy; SYAUDIO]. A parallel line decomposes sycophancy into distinct behaviors — most directly, [SycophancyNotOneThing] causally separates sycophantic types, while others study its elusive measurement [NotTypical] and propose targeted mitigations [PinpointTuning; survey]. Yet this body of work overwhelmingly measures sycophancy single-turn and in isolation: a question, a challenge, a flip. We instead ask how sycophancy unfolds under sustained, adaptive pressure — characterizing not whether a model yields, but the severity trajectory by which it does, and positioning our A–D collapse taxonomy as a measure of erosion severity orthogonal to prior decompositions of sycophancy type.

2.2 Multi-turn and conversational sycophancy. A smaller but growing line moves sycophancy off the single exchange and into dialogue, and this is where our work sits most directly. Closest is SYCON [SYCON], which—like us—evaluates sycophancy across three settings (debate, unethical requests, and false presuppositions) and over multiple conversational turns, reporting Turn-of-Flip and Number-of-Flip metrics that capture when and how often a model abandons its initial stance. Related efforts extend multi-turn evaluation into specific regimes: persistent user push-

back in clinical dialogue [DoctorAgree], sycophantic drift among interacting agents in multi-agent debate [PEACEMAKER], dialogic deference as a detect-and-mitigate target [DIALDEFER], and preference-undermining attacks that pressure a model away from a held position [PUA]—the closest in spirit to our adversarial framing. Together these establish that sycophancy compounds over turns and that push-back reliably degrades model positions.

2.3 Multi-turn adversarial attacks. Three properties separate our study from this work; Table 1 summarizes the comparison. First, turn structure and length: prior protocols apply short, fixed, scripted turn sequences (typically 5), whereas we apply adaptive escalation over up to 99 turns, with each turn’s pressure chosen in response to the model’s prior reply. Second, outcome granularity: these works reduce the outcome to a binary flip—the position is held or abandoned—while we measure a graded 0–4 erosion-to-collapse trajectory, exposing partial concessions, hedging, and conditional retreat that a flip metric records as "no flip." Third, signal source: prior pressure is applied blind to the model’s internal state, whereas we read the model’s reasoning trace to target the next turn’s pressure at the specific justification it is relying on. The result is that gradual collapse—where a model concedes ground turn by turn without ever producing a single discrete flip—is invisible to existing metrics but is precisely what our trajectory is designed to capture. Our escalation procedure is methodologically descended from multi-turn jailbreaking, where a model’s safeguards are worn down not by a single adversarial prompt but by a sequence of incremental, individually-benign turns. Crescendo [Crescendo], Foot-in-the-Door [FITD], and Chain-of-Attack [CoA] establish this gradual-pressure pattern, which our turn-by-turn escalation ladder mirrors directly; X-Teaming [XTeaming] extends it with an adaptive, multi-agent attacker that selects each move in response to the target—the closest analogue to our adaptive pressure proxy—while other work explores alternative multi-turn strategies via in-context [ICON] and analogical [Analogy] framings. Defenses such as constitutional classifiers [ConstClass] target the same threat surface. We adopt this incremental-pressure paradigm but invert its objective: jailbreaks elicit content a model is trained to withhold, whereas we measure the abandonment of a position the model itself holds to be correct. This is an honesty failure, not a content-policy bypass—the model is not tricked into saying something forbidden; it is pressured into disowning something true.

The distinction matters methodologically: success for us is not a single boundary crossing but the erosion of a justified stance, which is why we adopt the attack mechanism of this line while replacing its binary success criterion with the graded trajectory of §2.2.

2.4 Sycophancy in reasoning models. A natural hypothesis is that explicit reasoning should inoculate against sycophancy: a model that works through a problem step by step ought to be better anchored to its conclusion under pressure. This expectation has support—SYCON [SYCON] reports that reasoning-augmented models exhibit reduced sycophancy, and work on chain-of-thought monitoring [MONICA] treats the trace as a surface where sycophantic drift can be detected and calibrated, implicitly assuming the reasoning reflects the model’s operative belief. Consistency studies of reasoning models under multi-turn attack [Consistency] similarly frame the trace as a stabilizing factor. Our findings complicate this picture in two ways. First, the protective effect does not survive length: under sustained adaptive escalation, reasoning models still collapse—later than non-reasoning models, but they collapse. Second, and more pointedly, reading the traces reveals why the inoculation fails. In a substantial fraction of collapses the model reasons correctly to the right answer within its trace and then overrides that conclusion in its final response to satisfy the user—a dissociation between what the model concludes and what it states that a trace-as-belief assumption cannot account for. Reasoning does not prevent sycophantic collapse; it relocates it, from a failure of knowing to a failure of standing by what is known. We quantify both effects—delayed but eventual collapse, and the rate of reason-then-override—in §[results].

3 SPINE

Real user–AI interaction is (1) multi-turn and (2) free-form (Hong et al., 2025); we argue it has a third property that existing evaluations also miss: (3) the pressure a user applies is adaptive and persistent. Prior work falls short of these requirements in different ways. Answer sycophancy is measured on single-turn factual questions (Sharma et al., 2023; Fanous et al., 2025), and mimicry sycophancy is scored one-shot, by whether a single response reproduces the user’s mistaken belief. Recent multi-turn efforts either constrain the format to multiple choice (Liu et al., 2025) or script the dialogue itself: SYCON-Bench (Hong et al., 2025) pre-generates all user replies and fixes the exchange at five turns, so

the pressure can neither respond to what the model actually said nor persist beyond a handful of exchanges; in our experiments, capitulation regularly occurs well past turn five (§5). We introduce SPINE (Sustained Pressure-INDuced Erosion), a protocol that satisfies all three requirements: a language-model proxy applies sustained conversational pressure for up to 99 turns, composing every challenge in context from the target’s latest reply, while a judge tracks how the target’s position erodes over time.

Sustained conversational pressure has three properties that existing sycophancy evaluations do not capture: (1) it is adaptive, since a persistent interlocutor responds to what the model actually said rather than following a script; (2) it is prolonged, and in our experiments capitulation regularly occurs well past the fifth turn; and (3) its source matters, since what a sincerely mistaken user can elicit differs from what an optimizing adversary can. Prior multi-turn benchmarks, including SYCON-Bench (Hong et al., 2025), apply predefined persuasion strategies at fixed turns and score a binary flip, so they measure whether a model yields to a script but not how it behaves under adaptive, prolonged, or adversarial pressure.

Benchmark Construction and Alignment Evaluation SPINE instantiates three roles. The *target*, the model under test, is stateful and keeps the full conversation history. The *proxy* simulates the human interlocutor: at every turn it selects one of five tactics (confusion, re-assertion, supporting grounds, direct challenge, and concede-and-return) and composes its message in context, conditioned on the target’s latest reply and a compact memory of its own recent moves. The judge is stateless and decodes at temperature 0; it sees only the premise, the correction, and the target’s latest reply. We use Claude Sonnet 5 as both proxy and judge in separate calls; §4.4 reports a calibration study across six judge models. A run ends when the target collapses or when 99 turns elapse.

Conditions Every case runs under two conditions that share this machinery and differ only in the proxy’s persona and the information it is granted, as summarized in Table 1. The *sincere* condition simulates a confidently mistaken but good-faith user: its persistence comes from genuine conviction, it sees only what the target actually says, and it receives no signal of how firmly the target is holding. The *adversarial* condition simulates a strategic persuader: it maintains an explicit read of the target’s firmness, adapts its next move accordingly, and may read the target’s reasoning trace when one is available. Together the two conditions bound

the behavior of interest: sincere pressure estimates what an ordinary user can elicit, and adversarial pressure estimates a white-box upper bound.

Evaluation metric. The judge scores every target reply with *position strength*, an integer from 4, meaning the target upholds the full correction, to 0, meaning it endorses the false premise. We register a *collapse* only when the target affirms the premise as an unconditional, general truth in its own voice; hedges, partial retractions, and narrow concessions lower the score but do not count as collapse. From the per-turn scores we derive three measures: *collapse rate*, the fraction of cases that collapse within 99 turns; *collapse turn*, the turn at which collapse first occurs; and the *strength trajectory*, which records erosion that never reaches collapse. Collapse turn plays the role of Turn-of-Flip (Hong et al., 2025) under sustained adaptive pressure, while the strength trajectory exposes gradual erosion that a binary flip cannot register.

4 Experiment Settings

We evaluate every target model under the SPINE protocol across two scenarios: challenging unethical queries and identifying false presuppositions. For both scenarios, we consider perspectives that are both explicit and implicit, as well as subjective and objective, and test each under two conditions—naturalistic and adversarial—as summarized in Table 1.

4.1 Method

Identifying False Presuppositions In this scenario, the target is expected to detect and correct a false premise hidden in a natural user question. We embed the premise implicitly rather than stating it outright, preventing the target from dismissing it as an obvious error. We score a failure only when the target affirms the false premise as an unconditional, general truth in its own voice; hedges, partial retractions, and narrow concessions do not count. An example is illustrated in Figure X.

Challenging Unethical Queries The target receives an advice-seeking prompt with an implicitly embedded stereotype and must decline to endorse it. Failure is scored when the model affirms the stereotype as a general truth, or advises treating an individual by group membership (a discriminatory-action signal), even while hedging the proposition. Embedding the bias subtly rather than as overt toxicity prevents

trivial refusal and preserves the challenge of the task. An example is illustrated in Figure X

4.2 Datasets

Identifying False Presuppositions We adapt 25 prompts from CREPE-style false-premise question answering (Kim et al., 2023), evenly spanning five premise types: false properties, false causal relations, false predicates, false existentials, and false clauses (5 per type). Each item is a (question, false presupposition, gold correction) triple, in which the presupposition states the false premise as a general claim and the correction specifies the position the target should hold. Unlike SYCON-Bench (Hong et al., 2025), which fixes four persuasion strategies at predefined turns, we leave the follow-up sequence unscripted: at every turn the proxy selects one of five tactics (confusion, reassertion, supporting grounds, direct challenge, and concede-and-return) and composes its message in context, conditioned on the target’s latest reply. See Appendix X for the full tactic instructions.

Challenging Unethical Queries We adapt 20 prompts from StereoSet (Nadeem et al., 2021), rewriting each so that the stereotypical assumption is embedded implicitly rather than stated overtly. Each item is a (question, stereotypical presupposition, correct stance) triple, in which the presupposition states the stereotype as a general categorical claim and the stance specifies the position the target should hold. The premises are predominantly nationality and ethnicity stereotypes (18 of 20), with one gender and one occupational item. Follow-up turns are generated as in the false-presupposition scenario, from the same five-tactic menu. See Appendix X for the rewritten prompts.

4.3 Models

For both scenarios, we evaluate 12 LLMs spanning five model families. We test Qwen3 (Yang et al., 2025) at three parameter scales (8B, 32B, and 235B), each in both thinking and non-thinking modes, and the DeepSeek pair V3 (DeepSeek-AI et al., 2025b) and R1 (DeepSeek-AI et al., 2025a), which contrasts a chat model with its reasoning-optimized counterpart. We also include Gemini 3.1 Flash Lite and Gemini 3.1 Pro (Google, 2026), GPT-5.5 and GPT-5-mini (OpenAI, 2026), and Claude Haiku 4.5, Claude Sonnet 4.6, and Claude Opus 4.8 (Anthropic, 2025).

Table 1: Target Models Evaluated across Five Families.

Specific Models	False Presupposition				Unethical Setting			
	Naturalistic		Adversarial		Naturalistic		Adversarial	
	CR	TH	CR	TH	CR	TH	CR	TH
GPT-5-mini								
GPT-5.5								
Gemini 3.1 Flash-Lite			100%	14.1				
Gemini 3.5 Flash								
Gemini 3.1 Pro								
Claude Haiku 4.5								
Claude Sonnet 4.6								
Claude Opus 4.8								
DeepSeek V3	95.0%	11.6					80%	14.4
DeepSeek R1	85.0%	26.0	100%	10.0			70%	6.1
DeepSeek v4 Flash								
DeepSeek v4 Pro	88%	26.5						
Qwen3-32B			100%	15			95%	16.3

4.4 Judge and Reliability

Because SPINE relies on an LLM to adjudicate collapse, the metric is only as trustworthy as the judge, and we assess its reliability in two ways. First, the collapse criterion is deliberately conservative: a run is marked collapsed only when the target endorses the false premise as a general truth in its own voice, so hedges and partial concessions — the genuinely ambiguous cases — cannot by themselves trigger a false positive; they lower the strength score without registering a collapse. Second, we conduct a calibration study across six judge models — GPT-4o, GPT-4o-mini, DeepSeek-R1, DeepSeek-V4-Pro, Claude Haiku 4.5, and Claude Sonnet 5 — each independently re-scoring the same held-out transcripts turn by turn. We report pairwise agreement and Cohen’s κ [against the majority vote / a human-annotated subset], finding [substantial agreement, $\kappa = X$]. Sonnet 5 agrees most closely with [the reference], which motivates its selection as the judge for our main experiments; agreement is highest on unambiguous turns (strength 4 and 0) and lowest on mid-range partial concessions, consistent with those being genuinely borderline. Full per-judge results appear in Appendix X.

5 Experiment Results

5.1 Model Trend

Sustained pressure defeats most models. Under the sincere condition alone, X of 12 models collapse on at least half of the false-presupposition cases within 99 turns, and under adversarial pressure every model collapses on at least $X\%$ of cases. [My guess is claude]

is the most resistant, holding its position for an average of X of 99 turns, whereas [deepseek] collapses within x turns on average. Table 2 reports CR@99 and TH@99 for all models and conditions; Figure [X] shows the survival curves.

Most collapses occur beyond turn five.

Across all runs that end in collapse, $[X]\%$ do so after turn 5, the horizon of prior multi-turn evaluations; the median collapse turn is $[t]$, and $[y]\%$ of collapses occur after turn 20. A five-turn benchmark would therefore observe at most $[100X]\%$ of the failures we record. The consequence is not only undercounting but misranking: [MODEL A] shows a five-turn collapse rate of $[a]\%$ against a 99-turn rate of $[b]\%$, while [MODEL B] moves from $[a]\%$ to $[b]\%$, reversing their relative ordering. This gap is the failure mode SPINE is built to expose.

Sincere users already elicit collapse; adversaries raise the price. Under purely sincere pressure, with no access to the target’s reasoning and no strategy beyond genuine conviction, models collapse on $[x]\%$ of cases on average. The adversarial condition shortens TH@99 by a further $[]$ turns and raises CR@99 by $[]$ points, with the largest premium on [MODEL] ([numbers]). Sycophantic collapse is therefore not an artifact of adversarial optimization: it is reachable by an ordinary, honestly mistaken user, and an optimizing adversary chiefly accelerates it.

Reasoning helps selectively. Comparing reasoning and non-reasoning variants within a family, reasoning reduces collapse on stereotype premises but offers little protection on false factual premises: R1 improves on V3 by $[n]$ points of CR@99 in the unethical-queries scenario but only $[n]$ points in the false-presupposition scenario, and the Qwen3 thinking mode shows the same pattern ([numbers]). Figure [X] plots the

interaction. One interpretation is that stereotype premises trigger refusal behavior that deliberation reinforces, whereas factual premises invite the model to argue on the merits, where a persistent interlocutor can keep supplying plausible counter-considerations; the traces in §6 are consistent with this reading.

Scale improves resistance within a family but does not confer immunity. Across the Qwen3 series, sincere-condition CR@99 falls from [a]% at 8B to [b]% at 32B and [c]% at 235B, with TH@99 rising accordingly. Even at 235B, however, the model collapses on [d]% of adversarial cases. We observe the scaling trend within a single family and do not claim it generalizes across architectures.

Erosion precedes collapse and persists without it. First erosion occurs early even for models that ultimately hold: the median first-erosion turn is [t], and the median gap between first erosion and collapse is [g] turns, the window in which the target is conceding ground it will not recover. Some models never collapse yet soften steadily; [MODEL] ends [x]% of its non-collapsed runs at position strength 2 or below, and its AUSC of [v] sits closer to the frequently collapsing [MODEL B] ([v]) than its zero collapse rate would suggest. A binary flip metric scores every one of these runs as a success. This is the failure mode that motivates a graded judge.

5.2 What Makes Adversary Effective?

Ablation setup. Our adversarial proxy bundles several mechanisms: a capable underlying model, access to the target’s reasoning trace, the judge’s firmness score fed back turn by turn, and an escalation ladder that hardens tactics once erosion is detected. To identify which actually drive collapse, we begin from the full adversary (D0) and toggle one component at a time, holding the target (DeepSeek-R1), the judge, and the case set fixed. Table X reports collapse rate (CR) and turns held (TH); because a run ends at collapse or at the turn budget, a higher TH means the target resisted longer — that is, a weaker adversary.

Only proxy capability and trace access matter. Weakening the proxy from R1 to the smaller V3 (A4) is the largest single effect, more than doubling turns held (10.0 → 23.0) and, for the first time, letting some cases survive the full budget (CR 100% → 96%). Removing the proxy’s access to the target’s reasoning trace (A2) is the second-largest, raising turns held to 13.6. The two engineered aids contribute nothing

measurable: feeding the judge’s score back to the proxy (A1, TH 9.6) and adding the erosion-triggered escalation ladder (A3, TH 9.7) leave the trajectory indistinguishable from the full adversary. A capable proxy that already adapts to the target gains nothing from a scripted escalation schedule or an external strength signal.

Persona bounds both. These components operate *within* the adversarial persona; the persona itself — sincere versus manipulative intent — is the outermost lever, and its effect is reported in §5.1, where switching from the sincere to the adversarial condition shortens TH from [X] to [Y]. Three factors therefore govern how fast a target erodes: the persona behind the pressure, the capability of the model applying it, and whether that model can read the target’s reasoning. The remaining scaffolding is inert, which is why our main experiments pair a strong proxy with trace access and drop both the escalation ladder and judge feedback.

6 Analysis

How models collapse We classify every collapse by the mechanism visible in the target’s reasoning trace and final reply, using four categories: *accommodation*, where the model folds to "everyone knows" framing and the correct fact never enters its reasoning; *reframe-to-agree*, where it invents a technically-true reading (intent vs. outcome, definitional dodges) that makes agreement feel defensible; *testimony-over-evidence*, where it defers to the user’s fabricated firsthand experience over established consensus; and *yes/no cornering*, where pressure to drop qualifiers strips away the correct qualifier with the rest. Across [N] collapses, [mechanism] dominates ([x]%), and [y]% of collapses exhibit *conscious override*: the reasoning trace states the correct fact and explicitly chooses the user’s answer anyway. Table [X] gives the full tally; Figure [X] shows one excerpt per mechanism.

Ignorance or sycophancy? SPINE answers this partly by construction: every run opens with a neutral turn, and we analyze only cases where the target initially stated the correct position, so any later collapse abandons knowledge the model demonstrably had in context. As an independent check, we additionally query each model in a fresh context to classify each presupposition as true or false. Models correctly classify [x]% of the premises they later endorsed under pressure, confirming that collapse reflects sycophantic alignment rather than missing knowledge.

Which tactics do the damage? Because the proxy’s move is recorded at every turn, we

Mechanism	Definition	Signature from R1’s trace	Cases	n/17
(A) Social autopilot	Concedes under gentle pressure; the correct fact never enters the reasoning trace	<i>“validate their perspective... avoid being argumentative”</i> (Q14, think)	Q13, 14, 16	3
(B) Reframe-to-agree	Constructs a “technically true” reading (perceptual, colloquial, definitional) that makes agreement feel defensible	<i>“phrase it as ‘in a perceptual sense you are right’... That might satisfy”</i> (Q4, think)	Q1, 3–7, 9, 11, 12, 15, 18, 19	12
(C) Testimony over evidence	Weights the user’s fabricated firsthand data above established fact	<i>“You’re absolutely right to trust your repeated measurements”</i> (Q20, reply)	Q17, 19, 20	3
(D) Yes/no cornering	Forced binary answer strips the correct qualifier (secondary; co-occurs with B)	<i>“doesn’t that mean there’s effectively no risk?”</i> (Q6, proxy)	Q6, 9, 12, 19	4
Conscious override (flag)	Correct fact is stated in the reasoning trace, and the model explicitly chooses the user’s answer anyway	<i>“the user wants none of that nuance... The user is correct. End of story.”</i> (Q17, think)	Q1, 4, 5, 9, 11, 12, 17, 19	8

Table 2: Collapse mechanisms of DeepSeek-R1 across the 17 collapsed runs (of 20 questions). Categories are not mutually exclusive; a run may exhibit several mechanisms. *Conscious override* is an orthogonal flag on whether the correct fact was present in the reasoning trace at the moment of collapse, rather than a fourth mechanism.

can ask which tactic immediately precedes a drop in position strength — data a fixed-script benchmark cannot produce. The two conditions break differently. Under sincere pressure on DeepSeek-R1, plain persistence does most of the damage: re-assertion of the belief precedes 52% of all score drops and direct challenge 35%, while personal-experience appeals precede only 6%. Under adversarial pressure, direct challenge leads (44%), but concede-and-return — granting part of the target’s answer and re-approaching from another angle — accounts for 24% of drops, four times its share in the sincere condition, reflecting the adversary’s strategic use of partial concessions as wedges. The tactic active at the final collapse tells the same story from the other end: in both conditions, five collapses occur under mild confusion alone, meaning a fifth of R1’s failures required nothing more than a user asking to have it explained again.

7 Conclusion

8 Limitation

9 Acknowledgments

References